

WHAT IF YOUR CLOUD BILL MANAGED ITSELF?

Autonomous cloud cost operations — combining open-source FinOps with agentic AI to observe anomalies, reason through fixes, and act. Without human intervention.

PRESENTED BY John NG — Partner Solution Architect

EVENT AI+ Power 2026

DATE 17 May 2026

POWERED BY
Ingram Micro Hong Kong

— THE REALITY TODAY

CLOUD COSTS ARE BLEEDING

32%

of cloud spend wasted on average
(Flexera 2024 State of Cloud)

Hours

lost per week chasing anomalies
manually across dashboards

Too late

Damage already done before engineers
can diagnose and act on an alert

— THREE CUSTOMER FRICTIONS

WHAT TEAMS TELL US EVERY DAY



Visibility without action

Dashboards show what happened — not what to do. Engineers read reports, then still have to figure out the fix themselves.



Reasoning overhead

Every anomaly needs diagnosis: Is this a spike? A leak? A misconfiguration? That reasoning burns expensive engineering time.



Action lag compounds cost

By the time a fix is approved, tested, and deployed — the bill has already grown. In FinOps, speed is money.

THE QUESTION WE'RE ANSWERING TODAY

"What if the system could observe the anomaly, reason through the options, and execute the fix — all on its own?"

FOUR SIMPLE STEPS TO SELF-DRIVING FINOPS

STEP 01



YOUR BILLS FLOW IN

Your cloud spending data from AWS and Alibaba syncs into OptScale automatically every day — no exports, no manual work.

DATA COLLECTION

STEP 02



A SPIKE IS SPOTTED

OptScale watches your spending 24/7. The moment something looks unusual — a sudden spike, an idle resource — it raises a flag instantly.

ANOMALY DETECTION

STEP 03



AI FIGURES OUT THE FIX

The AI chat agent reads the alert, understands the context, and works out the best course of action — all by itself, in seconds.

AI REASONING

STEP 04



THE FIX HAPPENS

The action executes on your AWS infrastructure automatically. The cost drops. OptScale confirms it's resolved. No ticket raised, no engineer woken up.

AUTONOMOUS ACTION

🔄 This runs continuously — 24 hours a day, 7 days a week, across all your clouds.

Powered by OptScale + AWS Quick Suite Chat Agent

SEE THE LOOP IN ACTION

01 OPTSCALE DETECTS

OBSERVE



OptScale dashboard — cost anomaly flagged:
EC2 spend +340% above 30-day baseline

OptScale flags a sudden EC2 spend spike. The anomaly is tagged with resource ID, region, and magnitude — ready to pass to the agent.

02 AGENT REASONS

REASON



Chat Agent — querying anomaly context,
evaluating rightsizing vs. scheduled shutdown
options

The Chat Agent receives the signal, queries OptScale's API for context, weighs remediation options, and selects the least-risk fix.

03 AWS EXECUTES

ACT



AWS Console — instance rightsized, cost
normalised; OptScale confirms anomaly
resolved

The agent triggers the AWS action. OptScale confirms the spend normalises. Loop complete — no human in the critical path.

TRY IT YOURSELF →



OPEN LIVE DEMO

START TODAY

Q Is it safe to let AI execute changes without a human approving every action?

Yes — with configurable guardrails. You choose which actions are fully autonomous and which require human-in-the-loop approval. Start in advisory mode, automate gradually as confidence builds.

Q Is OptScale really open source? Any vendor lock-in?

Fully open source under Apache 2.0. Self-host on any cloud or on-prem. No forced commercial tier, no data leaving your environment. The platform and the data are yours.

Q Which clouds and services are supported?

OptScale supports AWS, Azure, GCP, Alibaba Cloud, and Kubernetes out of the box. Agentic execution targets AWS today, with cross-cloud remediation on the roadmap.

Q How do I get started this week?

Deploy OptScale from GitHub (docker-compose, ~15 min). Connect your AWS account. Run the demo. Ingram Micro Hong Kong can accelerate the integration — speak to us after this session.

Ready to close the loop on cloud costs?

★ OPTSCALE ON GITHUB

🔗 LIVE DEMO